

Sieve: Training-Free, Model-Free, Question-Conditioned Context Compression for Efficient Long-Context LLM Inference: An Accuracy–Compression–Cost Frontier Study of Model-Free Information Selection

Chang Tan

Undergraduate Student
School of Economics and Management
Dalian University of Technology
No. 2 Linggong Road, Ganjingzi District
Dalian, Liaoning 116024
China

Corresponding author: Chang Tan
E-mail: 2352935097@qq.com

Manuscript type: Full Length Article.

Sole authorship: Chang Tan is the sole author of this manuscript and, as the only author, also serves as the corresponding author.

Funding: None.

Acknowledgements: None.

Declaration of competing interest: The author declares no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

CRedit authorship contribution statement: **Chang Tan:** Conceptualization, Methodology, Software, Investigation, Data curation, Writing – original draft, Writing – review & editing.

Data availability: The benchmarks are public (QASPER, HotpotQA, 2WikiMultihopQA). The released repository contains the pool-construction scripts with fixed seeds, the sampled pools, every raw checkpoint, and one command per table and figure, so that every number in this paper can be regenerated from scratch.